

# *High Performance Programming, Lecture 3*

*More programming in C*

## *Follow up on Lab 1*

Check your quota on unix systems using `quota -v`:

```
vranx> quota -v
Disk quotas for anakr123 (uid 340973):
Filesystem      usage  quota  limit   .....
/home/anakr123
                24889 1048576 1048576   .....

vranx>
```

Find largest files and directories:

```
du -a | sort -n
```

## *C building process*

4 steps: preprocessing, compilation, assembly, linking.

Look at C\_build\_process.pdf (see under Extra files in Student Portal) for more information.

Object files (.o) are “almost executables”. Contain: the machine code + metadata about the addresses of all used variables and functions (also called symbols).

*Undefined reference error:*

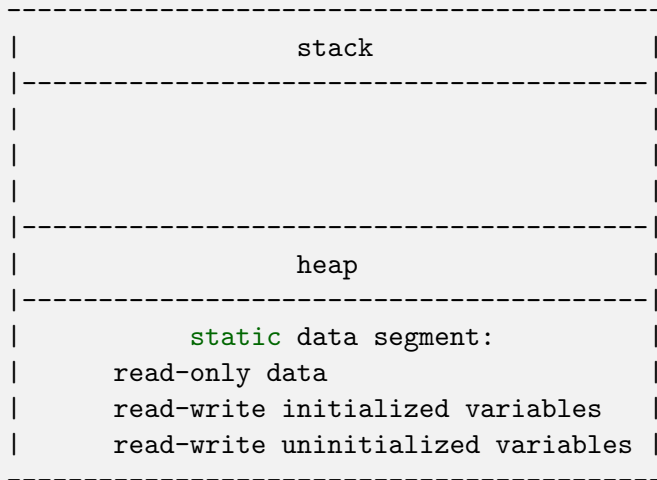
That means that this symbol is referenced or used in this object file, but its value was not defined here.

Undefined symbols cannot be used unless a **definition** is linked in via a library or object file!

## *Follow up on Lab 1*

### *Program's address space*

Running program's view of memory in the system:



## *Follow up on Lab 1*

### *Program's address space*

The stack is a LIFO (Last-In-First-Out) data structure. Memory allocation/deallocation is done automatically.

If memory is allocated on the heap, it exists as long as the developer wants.




Image: <https://stackoverflow.com>

# Follow up on Lab 1

*Program's address space (Try it!)*

Example program: test stack and heap

```
stack at 0x7fff0df1b697, heap at 0x1e12260
stack at 0x7fff0df1b667, heap at 0x1e12690
stack at 0x7fff0df1b637, heap at 0x1e126b0
stack at 0x7fff0df1b607, heap at 0x1e126d0
stack at 0x7fff0df1b5d7, heap at 0x1e126f0
```

```
#include <stdlib.h>
#include <stdio.h>
void f(int depth) {
    if (depth <= 0) return;
    char c;
    char *ptr = malloc(1);
    printf("stack at %p, heap at %p\n", &c, ptr);
    f(depth-1);}
int main() {
    f(5);
    return 0;}
```

*Note the difference!*

```
char *s1 = "Hello";
```

Put the literal string "Hello" in the read-only memory and let s1 point to it. Any writing operation on this memory illegal!

```
char s2[] = "Hello";
```

Put the literal string "Hello" in the read-only memory and copy it to newly allocated memory on the stack (s2 is an array of 6 characters: 'H', 'e', 'l', 'l', 'o', and '\0').  
Now `s[0] = 'J';` is legal.

## Follow up on Lab 1

### Compare

```
void assign_string(char
    *g)
{
    g = "Good!";
}
int main()
{
    char *s = "Hello";
    assign_string(s);
    printf("%s\n", s); //
    ‘‘Hello’’
    return 0;
}
```

```
void assign_string(char
    **g)
{
    *g = "Good!";
}
int main()
{
    char *s = "Hello";
    assign_string(&s);
    printf("%s\n", s); //
    ‘‘Good!’’
    return 0;
}
```



## *Follow up on Lab 1*

Consider the following code:

```
int * p = (int *) malloc (sizeof(int));
*p = 0; // !!!!!
for(int i = 0; i < N; i++) *p += i;
printf("%d\n", *p);
```

## *(Pseudo)random numbers*

rand returns a pseudo-random number in the range of 0 to RAND\_MAX.

srand sets seed for the pseudo-random number generator used by rand().

time returns the time since the Epoch (00:00:00 UTC, January 1, 1970) in seconds.

```

srand(time(NULL));    // Initialization (called
once).
int r = rand();        // Returns a pseudo-random
integer between 0 and RAND_MAX.
}

```

## *Follow up on Lab 1*

*Structure datatype. Compare:*

```
struct node {
    int val;
    struct node * next;
};
typedef struct node
    node_t;

node_t A;
```

```
typedef struct node {
    int val;
    struct node * next;
} node_t;

node_t A;
```

## *Followup after Lab 1*

*Small example programs available*

For those who are new to the C language, a few example programs are available in the Student Portal:

Documents → Example programs

- `input_args.c`: **argc** and **argv** in `main()` function
- `pointers.c`: playing around with **pointers**
- `file_io.tar.gz`: **reading and writing files** using `fopen()`, `fread()`, `fwrite()`, etc. Reading and writing text files.
- `struct_example.c`: simple example of using a **struct** datatype
- `func_with_ptr_arg.c`: simple examples of using **pointer arguments** to functions
- `read_input_data_into_array.c`: reading input characters into an array
- and some more

## *If you missed Lab 1*

Look at the document CONTENT → “Getting started with linux” in the Student Portal.

### **Labs are mandatory.**

The rule is:

if you miss a lab you must submit a small report for that lab instead, no later than **3 working days after the lab**.

Submit using Student portal under ASSIGNMENTS → “Extra reports” or send me by email (first option is preferable)

## *Follow up on Lab 1*

Any questions?

Check that they get marked as completed in the Progress function in the Student Portal.

# Assignment 1

## IMPORTANT NOTES:

- no binary files such as object files or executable files should be included
- makefile should have the name “Makefile”, use the checking script! It should finish with “Congratulations” message.
- specify dependencies correctly, if file is changed it should be recompiled
- read carefully “Preparing your submission” part!

Submit using Student portal. Deadline is this Friday (February 1)!

## Assignment 2

Assignment 2 is available in Student Portal. Submit your solution using Student portal. Deadline is next Friday (February 8)!

For Assignments 2-6 you are expected to **work in pairs**, “groups” with two students/group, using the “group division” function in the Student Portal. When you are in a group, you can access the files for Assignment 2.

If you **do not have a teammate, send email** to Anastasia about it, then I can pair you up with someone.



Today:

- Linked list (very briefly)
- Multidimensional arrays
- Pointers to functions
- Timing and complexity (depends on time left)

Next time:

Optimization – Part I (how to do less work)

## *Linked list*


Linked list is a simple example of a dynamic structure which uses pointers.

**Advantages** over arrays:

- Items can be added or removed from the middle of the list;
- No need to define an initial size.

**Disadvantages:**

- No random access;
- Uses pointers  $\implies$  complications of code;
- Larger overhead due to dynamic allocation.

*Linked list*

Define a linked list node:

```

struct node {
    int val;
    struct node * next;
};

typedef struct node node_t;

```

## *Linked list*

Add item to the end of the list:

```

void push(node_t * head, int val) {
    node_t * current = head;
    while (current->next != NULL) {
        current = current->next;
    }

    current->next = (node_t *) malloc(sizeof(node_t));
    current->next->val = val;
    current->next->next = NULL;
}

```

## *Linked list*

Look at the tutorial page:

[https://www.learn-c.org/en/Linked\\_lists](https://www.learn-c.org/en/Linked_lists).

You will need to create your own linked list in Assignment 2.

## Multidimensional arrays

General form: `type name[size1][size2]...[sizeN];`

Arrays are always laid out *contiguously* in memory.

2D array initialization:

```
int arr[2][5] = {  
    {0, 1, 2, 3} ,  
    {5, 6, 7, 8, 9}  
};  
int value= a[1][3]; // value = 8
```

## Multidimensional arrays

General form: `type name[size1][size2]...[sizeN];`

Arrays are always laid out *contiguously* in memory.

2D array initialization:

```
int arr[2][5] = {  
    {0, 1, 2, 3} ,  
    {5, 6, 7, 8, 9}  
};  
int value= a[1][3]; // value = 8
```

`arr[2][5]` is a matrix with 2 rows and 5 columns (array of arrays)

In memory it is stored in row-major order: after `arr[0][3]` comes `arr[0][4]` and after `arr[0][4]` comes `arr[1][0]`.

0	1	2	3	0	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

## Multidimensional arrays

2D array initialization:

```
int arr[2][5] = {  
    {0, 1, 2, 3} ,  
    {5, 6, 7, 8, 9}  
};
```

*//or*

```
int arr[2][5] = { 0, 1, 2, 3, 0, 5, 6, 7, 8, 9};
```

*// or skipping the first dimension (not the second!)*

```
int arr[][5] = {  
    {0, 1, 2, 3} ,  
    {5, 6, 7, 8, 9}  
};
```

0	1	2	3	0	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---



## *Dynamic memory allocation*

Allocate memory for a matrix  $3 \times 5$  dynamically.

```
int **arr = (int **)malloc(3 * sizeof(int*));
for (i = 0; i < 3; i++)
    arr[i] = (int *)malloc(5 * sizeof(int));
```

## *Dynamic memory allocation*

Allocate memory for a matrix  $3 \times 5$  dynamically.

```
int **arr = (int **)malloc(3 * sizeof(int*));
for (i = 0; i < 3; i++)
    arr[i] = (int *)malloc(5 * sizeof(int));
```

How to free the allocated memory?

## *Dynamic memory allocation*

Allocate memory for a matrix  $3 \times 5$  dynamically.

```
int **arr = (int **)malloc(3 * sizeof(int*));
for (i = 0; i < 3; i++)
    arr[i] = (int *)malloc(5 * sizeof(int));
```

How to free the allocated memory?

```
for (int i = 0; i < 3; i++)
    free(arr[i]);
free(arr);
```

## Dynamic and static memory allocation of a matrix

### Compare:

```
int arr[2][3]; // 2D array, size is known at the
               compile time
int* arr[3];   // array of pointers
int **arr;     // pointer to pointer
```

### Example:

```
int main(int argc, char *argv[]) { /* ... */ }
// or
int main(int argc, char **argv) { /* ... */ }
```

## *Pointers to functions*

Declaration of a function:

```
function_return_type function_name(function argument
    list)
```

Declaration of a pointer to a function:

```
function_return_type (*pointer_name)(function
    argument list)
```

For example (note parentheses () here):

```
void    (*ptrfun)()
void    (*ptrfun)(double, char)
double  (*ptrfun)(double, char *)
int*    (*ptrfun)(int*)
```

Check out this “dictionary” for C/English: <https://cdecl.org/>.

## Pointers to functions

```

#include <stdio.h>
#include <stdlib.h>
double add(double a, double b)
{return a+b;}

void print_output(double a, double b, double (*funptr)(
    double, double))
{printf("Value is %lf \n", (*funptr)(a, b));}

int main(int argc, char const *argv[])
{
    int n = atoi(argv[1]);
    double (*funptr)(double, double); // declare pointer to
        function
    funptr = &add; // assign address of a function
    print_output(3, 4, funptr);
    return 0;
}

```

*qsort*

qsort function (from stdlib.h) for sorting arrays:

```
void qsort (void* base, size_t num, size_t size, int
            (*compar)(const void*,const void*));
```

For types that can be compared using regular relational operators, a general compar function may look like:

```
int compareMyType (const void * a, const void * b)
{
    if ( *(MyType*)a <  *(MyType*)b ) return -1;
    if ( *(MyType*)a == *(MyType*)b ) return 0;
    if ( *(MyType*)a >  *(MyType*)b ) return 1;
}
```

Note: for strings we should use strcmp!

*qsort (Test it!)*

```

#include <stdio.h>
#include <stdlib.h>

int CmpDouble(const void * p1, const void *p2)
{
    double a = *(double *)p1;
    double b = *(double *)p2;
    if(a > b) return -1;
    if(a < b) return 1;
    return 0;
}

int main() {
    double arrDouble[] = {9.3, -2.3, 1.2, -0.4, 2, 9.2, 1, 2, 0};
    int arrDoubleLen = sizeof(arrDouble) / sizeof(double);
    qsort (arrDouble, arrDoubleLen, sizeof(double), CmpDouble);
    for (int i=0; i<arrDoubleLen; ++i)
        printf("%d: %lf\n", i, arrDouble[i]);
    return 0;
}

```



## *Another example (later in the course)*

Function for creating a new POSIX thread:

```
#include <pthread.h>

int
pthread_create (pthread_t *thread_id,
                const pthread_attr_t *attributes,
                void *(*thread_function)(void *),
                void *arguments);
```

thread\_function is the function the new thread is executing.

## *Time measuring*


Bash command `time`.


C commands:

- `gettimeofday()`
- `clock_gettime()` on Solaris or Linux, or `clock_get_time()` on Mac.

## Time measuring




 = user time

 = system time

 = some other user's time

 +  +  = real (wall clock) time

*Usually the word “time” refers to user time.*

 cumulative user time

## *The time command*

Type in your terminal:

```
$ time ./executable
```

You will get something like this:

```
$ time ./executable
real          0m0.143s
user          0m0.001s
sys           0m0.010s
$
```

The time command gives three timing measurements:

- **real**: “wall-clock time”
- **user**: time spent in user code, i.e. what you’ve written
- **sys**: time spent in system calls, e.g. malloc, printf, etc

## *The time command*

Unreliable if runtime is very short.

→ *Make sure to run your code long enough!*

user + sys will tell you how much actual CPU time your process used. Can potentially exceed the wall clock time in multi-threaded programs.

## Time measuring

```
int gettimeofday(struct timeval *tv, struct timezone
                 *tz);
```

gives the number of seconds and microseconds since the Epoch  
(1970-01-01 00:00:00 (UTC))

```
struct timeval {
    time_t      tv_sec;    /* seconds */
    suseconds_t tv_usec;   /* microseconds */
};

struct timezone {
    int tz_minuteswest; /* minutes west of Greenwich */
    int tz_dsttime;     /* type of DST correction */
};
```

## *Time measuring gettimeofday*

### **In your code:**

Do not forget to `#include <sys/time.h>`  
(You may need to link with `-lrt`)

```
struct timeval t0, t1;
gettimeofday(&t0, 0);
/* your code */
gettimeofday(&t1, 0);

long elapsed_time_usec = (t1.tv_sec-t0.tv_sec)*1e6 + t1.
    tv_usec-t0.tv_usec;
double elapsed_time_sec= (t1.tv_sec-t0.tv_sec) + (t1.
    tv_usec-t0.tv_usec)/1e6;

printf("%ld microsec, %lf sec\n", elapsed_time_usec,
    elapsed_time_sec);
```

## Time measuring

```
int clock_gettime(clockid_t clk_id, struct timespec *
    tp);
```

gives the number of seconds and **nanoseconds** since the Epoch (1970-01-01 00:00:00 (UTC))

`clk_id` is `CLOCK_REALTIME` *// we will use this one*

```
struct timespec {
    time_t    tv_sec;           /* seconds */
    long      tv_nsec;         /* nanoseconds */
};
```



## *Time measuring clock\_gettime*

### **In your code:**

Do not forget to `#include <time.h>`  
(You may need to link with `-lrt`)

```
struct timespec t0, t1;
clock_gettime(CLOCK_REALTIME, &t0);
/* your code */
clock_gettime(CLOCK_REALTIME, &t1);

long elapsed_time_nsec = (t1.tv_sec-t0.tv_sec)*1e9 + t1.
    tv_nsec-t0.tv_nsec;
double elapsed_time_sec = (t1.tv_sec-t0.tv_sec) + (t1.
    tv_nsec-t0.tv_nsec)/1e9;

printf("%ld nano sec, %lf sec\n", elapsed_time_nsec,
    elapsed_time_sec);
```

## *Guidelines for measuring execution time*

- Use a high resolution timer, such as `clock_gettime()`
- Check the amount of time your program spends in the OS using the `time` command
- Do not trust too short timings
- Measure several runs to calibrate variability (3 or more)
- Pick a shortest or an average time as a representative
- Outliers can be important!

## *Example*

Example program in the Student Portal:

Documents → Example programs → `timings.c`: measuring **timings** for different parts of a program using `gettimeofday()`

## Complexity

The same problem often can be solved by different algorithms.

*Which algorithm to choose?*

**Computational complexity:** how many resources we need in order to solve some problem?

We want to compare algorithms, not computers!

## *Complexity*

**Space complexity** - memory needed for an algorithm to solve a given problem. We are measuring total allocated memory in some units.

**Time complexity** - time needed for an algorithm to solve a given problem. Time is measured in some units, for example seconds or minutes, it can be number of cycles.

## Complexity

### Examples:

- $n! = 1 * 2 * 3 \dots (n-1)n$ :  $n-1$  multiplications.
- nested loops:  $n \times m$  function calls

```
for(int i = 0; i < n; ++i)
    for(int j = 0; j < m; ++j){ f(); }
```

- matrix multiplication:  $2n^2$  storage,  $n^2(2n-1) = 2n^3 - n^2$  operations.

The time required to solve each of these problems will depend on a computer and on an implementation. With increasing the problem size  $n$  the time will in general increase.

We consider just a **dominant part** of the instruction count. Matrix multiplication requires  $\approx 2n^3$  operations.

## Complexity

Matrix multiplication of matrices of size  $n$  on a computer X takes  $30n^3$  microseconds:

$n = 100$ , it needs 30 seconds

$n = 200$ , it needs 240 seconds - **8 time more!**

On another computer Y multiplication takes  $0.3n^3$  microseconds:

$n = 100$ , it needs 0.3 seconds

$n = 200$ , it needs 2,4 seconds - **8 time more!**

**We want to compare algorithms, not computers!**

Implementations on similar computer architecture may give different timings up to a constant.

Matrix multiplication requires  $cn^3$  operations, where  $c = \text{const.}$

## Complexity


Introduce function  $f$  which gives a feeling about the amount of work required for a given problem size,  $f$  is monotonously growing.

Let  $f$  and  $g$  are functions from  $S \subset \mathbb{R}$  to  $\mathbb{R}$ .

$f$  is not growing faster than  $g$  if

$\exists x_0 \in S$  and  $c > 0$  such that

$\forall x > x_0, |f(x)| < c|g(x)|$ .



We denote such relation as  $\mathbf{f} \in \mathbf{O(g)}$  (when  $x \rightarrow \infty$ ) — it says that the algorithm has an **order g** complexity.



# Complexity

Examples:

$$-2x^3 + 4x^2 + x = O(x^3)$$

$$n/2 = O(n)$$


$$\log n + n - 2 = O(n)$$

## Complexity

Common time complexities:

- $\mathcal{O}(1)$ , i.e. constant  $\leftarrow$  access element in an array
- $\mathcal{O}(\log n)$ , i.e. logarithmic  $\leftarrow$  search element in a binary search tree
- polynomial:
  - $\mathcal{O}(n)$  linear  $\leftarrow$  compute  $n!$
  - $\mathcal{O}(n \log n)$   $\leftarrow$  quick sort
  - $\mathcal{O}(n^2)$   $\leftarrow$  matrix addition
  - $\mathcal{O}(n^3)$   $\leftarrow$  matrix multiplication
- $\mathcal{O}(2^{\text{poly}(n)})$ , i.e. exponential  $\leftarrow$  Fibonacci numbers  $\mathcal{O}(2^n)$
- $\mathcal{O}(n!)$ , i.e. factorial  $\leftarrow$  find all permutations of a string

# Complexity



## *n-body problem*

Example problem: evolution of the galaxy (N stars) in a selected space region.



Simple approach: each star acts with some force to other N-1 stars. Then the total number of force calculations is  $\mathcal{O}(N^2)$ . If the number of particles is increased by factor ten the number of calculations increases by a factor of 100...

Use approximation! A tree based approximation scheme reduces the computational complexity of the problem from  $\mathcal{O}(N^2)$  to  $\mathcal{O}(N \log N)$  (The Barnes-Hut algorithm).

The basic idea: if the group is sufficiently far away, we can approximate its gravitational effects by using its center of mass.

That's all